

Submission Assignment #2

Coordinator: Jakub Tomczak

Name: Marton Fejer, Netid: [2694002]

1 Question answers

Question 1 Let $f(X, Y) = X / Y$ for two matrices X and Y (where the division is element-wise). Derive the backward for X and for Y . Show the derivation.

For the backward functions we use the chain rule to get the following:

$$\begin{aligned} \text{Backward}(X) &= \frac{\partial \text{Loss}}{\partial f(X, Y)} \cdot \frac{\partial f(X, Y)}{\partial X} \\ \text{Backward}(Y) &= \frac{\partial \text{Loss}}{\partial f(X, Y)} \cdot \frac{\partial f(X, Y)}{\partial Y} \end{aligned}$$

Figure 1: equation (1)

Thus, we can calculate the partial derivatives using the quotient rule to obtain the backward functions for X and Y :

$$\frac{\partial f(X, Y)}{\partial X} = \frac{\partial}{\partial X} \left(\frac{X_{ij}}{Y_{ij}} \right) = \frac{1 \cdot Y_{ij} - X_{ij} \cdot 0}{Y_{ij}^2} = \frac{Y_{ij}}{Y_{ij}^2} = \frac{1}{Y_{ij}}$$

Figure 2: equation (2)

$$\frac{\partial f(X, Y)}{\partial Y} = \frac{\partial}{\partial Y} \left(\frac{X_{ij}}{Y_{ij}} \right) = \frac{0 \cdot Y_{ij} - X_{ij} \cdot 1}{Y_{ij}^2} = \frac{-X_{ij}}{Y_{ij}^2}$$

Figure 3: equation (3)

Question 2 Let f be a scalar-to-scalar function. Let $F(X)$ be a tensor-to-tensor function that applies f element-wise (For a concrete example think of the sigmoid function from the lectures). Show that whatever f is, the backward of F is the element-wise application of f' applied to the elements of X , multiplied (element-wise) by the gradient of the loss with respect to the outputs.

For the backward of $f(x)$, we have:

$$\frac{\partial \text{Loss}}{\partial f(x)} \cdot \frac{\partial f(x)}{\partial x}$$

Figure 4: equation (4)

then, for the backward of $F(x)$, we get the following:

Thus, the backward of $F(x)$ would still only be an element-wise function.

$$\frac{\partial Loss}{\partial F(x_{ijk})} \cdot \frac{\partial F(x_{ijk})}{\partial x}$$

Figure 5: equation (5)

Question 3 Let matrix W be the weights of an MLP layer with f input nodes and m output nodes, with no bias and no nonlinearity, and let X be an n -by- f batch of n inputs with f features each. Which matrix operation computes the layer outputs? Work out the backward for this operation, providing gradients for both W and X .

For clarity, the number of features "f" will be referred to as "g" instead. The matrix of the weights has the shape of $(f \times m)$. The matrix describing X (so the input) has the shape of $(n \times g)$. The layer outputs must have the form of $(m \times g)$, thus in order to obtain that form, one must calculate the multiplication of the transpose of the Weights together with the Input matrix:

$$\text{Layer outputs} = W^T X$$

Figure 6: Matrix operation to calculate the layer outputs.

Thus, the backward for this operation wrt to W or X is the following:

$$\begin{aligned} \frac{\partial Loss}{\partial W} &= \frac{\partial Loss}{\partial \text{Layer Outputs}} \cdot \frac{\partial \text{Layer Outputs}}{\partial W} \\ \frac{\partial Loss}{\partial X} &= \frac{\partial Loss}{\partial \text{Layer Outputs}} \cdot \frac{\partial \text{Layer Outputs}}{\partial X} \end{aligned}$$

Figure 7: The backward of layer output wrt W and X .

whereby the gradient of the layer output wrt W and X is the following:

$$\begin{aligned} \frac{\partial \text{Layer Outputs}}{\partial W} &= \frac{\partial}{\partial W} (W^T X) = X \\ \frac{\partial \text{Layer Outputs}}{\partial X} &= \frac{\partial}{\partial X} (W^T X) = W^T \end{aligned}$$

Figure 8: Gradients of the layer output wrt W and X .

Question 4 Let $f(x) = Y$ be a function that takes a vector x , and returns the matrix Y consisting of 16 columns that are all equal to x . Work out the backward of f .

The function $f(x)$ takes the vector x of shape $(n \times 1)$ and turns it into matrix Y of shape $(n \times 16)$. Thus, the function $f(x)$ multiplies the vector x with another vector z of shape (1×16) :

$$f(x) = x \cdot z$$

Figure 9: Function $f(x)$.

and thus, the backward of $f(x)$ wrt to x is the following:

$$\begin{aligned} \frac{\partial Loss}{\partial x} &= \frac{\partial Loss}{\partial f(x)} \cdot \frac{\partial f(x)}{\partial x} \\ &= \frac{\partial Loss}{\partial f(x)} \cdot \frac{\partial}{\partial x} (x \cdot z) = \frac{\partial Loss}{\partial f(x)} \cdot z \end{aligned}$$

Figure 10: Backward of the function $f(x)$ wrt x .

Question 5

- 1) What does `c.value` contain?
- 2) What does `c.source` refer to?
- 3) What does `c.source.inputs[0].value` refer to?
- 4) What does `a.grad` refer to? What is its current value?

- 1) `c.value` contains the outcome of the tensor addition, so if `a` and `b` are tensors of the same shape, then `c.value` is the element-wise addition of `a` and `b`.
- 2) `c.source` refers to the operation node in the sense that it refers to what operation (together with which inputs) resulted in the node `c`. `c.source` is itself an operation node object.
- 3) `c.source.inputs[0].value` refers to the value(s) of the tensor node `a`. `c.source.inputs[0]` itself is a tensor node.
- 4) `a.grad` refers to the gradient of the loss wrt to `a.value`. Its current value is 0.

Question 6

You will find the implementation of `TensorNode` and `OpNode` in the file `vugrad/core.py`. Read the code and answer the following questions:

- 1) An `OpNode` is defined by its inputs, its outputs and the specific operation it represents (i.e. summation, multiplication). What kind of object defines this operation?
- 2) In the computation graph of question 5, we ultimately added one numpy array to another (albeit wrapped in a lot of other code). In which line of code is the actual addition performed?
- 3) When an `OpNode` is created, its inputs are immediately set, together with a reference to the op that is being computed. The pointer to the output node(s) is left `None` at first. Why is this? In which line is the `OpNode` connected to the output nodes?

- 1) An `OpNode` performs some operation on some inputs (a tensor). The result of this operation is the output. The operation itself is a class object.
- 2) The addition is done with a function in `core.py` at line 104-108. This leads to the actual operation class "Add" which performs the actual addition at line 324.
- 3) The reason it's left as `None` at first, is that it will depend on whether one does a forward or a backward pass with it (in which case they will differ). The `OpNode` is connected to the outputs at line 248 and 249 (first defined at 248, then connected at 249) in `core.py`.

Question 7

When we have a complete computation graph, resulting in a `TensorNode` called `loss`, containing a single scalar value, we start backpropagation by calling

`loss.backward()`

Ultimately, this leads to the `backward()` functions of the relevant Ops being called, which do the actual computation. In which line of the code does this happen?

This happens in line 97 (see fig. 11) whereby the "relevant operation" is `self.source` and the `backward()` function call is its backward implementation.

```

94 | # If we've been visited by all parents, move down the tree
95 | if self.visits == self.numparents or start:
96 |     if self.source is not None:
97 |         self.source.backward()
```

Figure 11: Backward of the function $f(x)$ wrt x .

Question 8

core.py contains the three main Ops, with some more provided in *ops.py*. Choose one of the ops *Normalize*, *Expand*, *Select*, *Squeeze* or *Unsqueeze*, and show that the implementation is correct. That is, for the given forward, derive the backward, and show that it matches what is implemented.

For normalization, the backward function has the general form:

$$\text{Backward}(N(x_i)) = \frac{\partial \text{Loss}}{\partial N(x_i)} \cdot \frac{\partial N(x_i)}{\partial x_i}$$

Figure 12: Backward of the Loss wrt x_i .

the derivative calculation is as follows:

$$\begin{aligned} \frac{\partial N(x_i)}{\partial x_i} &= \frac{\partial}{\partial x_i} \left(\frac{x_i}{\sum_j^n x_j} \right) \\ &= \frac{1 \cdot \sum_j^n x_j - x_i}{\left(\sum_j^n x_j \right)^2} \\ &= \frac{1}{\sum_j^n x_j} - \frac{x_i}{\left(\sum_j^n x_j \right)^2} \end{aligned}$$

Figure 13: Derivative of the function $N(x_i)$ wrt x_i .

Replacing the local derivative in the backward function with the calculated derivative, one gets:

$$= \frac{\frac{\partial \text{Loss}}{\partial N(x_i)}}{\sum_j^n x_j} - \frac{\frac{\partial \text{Loss}}{\partial N(x_i)} \cdot x_i}{\left(\sum_j^n x_j \right)^2}$$

Figure 14: Complete backward function of the Loss wrt x_i .

Comparing it to the backward function of normalization in *ops.py*, one can indeed see that it is the same:

```

127 class Normalize(Op):
128     """
129     Op that normalizes a matrix along the rows
130     """
131     @staticmethod
132     def forward(context, x):
133
134         sumd = x.sum(axis=1, keepdims=True)
135
136         context['x'], context['sumd'] = x, sumd
137
138         return x / sumd
139
140     @staticmethod
141     def backward(context, go):
142
143         x, sumd = context['x'], context['sumd']
144
145         return (go / sumd) - ((go * x)/(sumd * sumd)).sum(axis=1, keepdims=True)
146

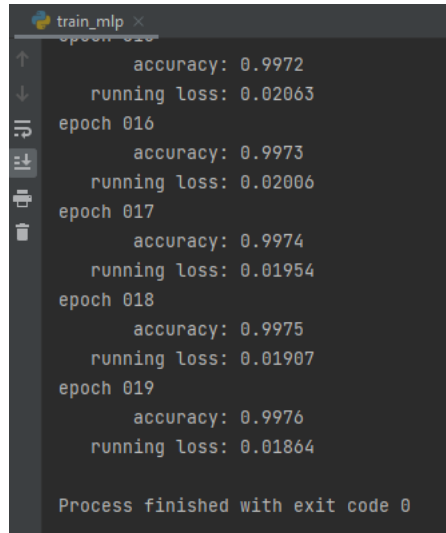
```

Figure 15: Backward function for normalization implemented in *ops.py*

Question 9

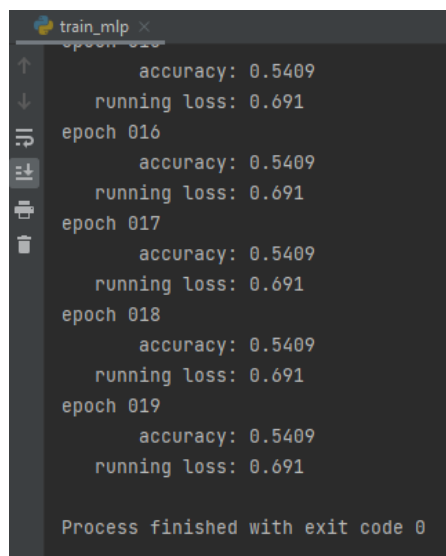
The current network uses a Sigmoid nonlinearity on the hidden layer. Create an Op for a ReLU nonlinearity (details in the last part of the lecture). Retrain the network. Compare the validation accuracy of the Sigmoid and the ReLU versions.

Retraining the network with Relu non-linearity yields a lower accuracy and higher loss for the simple synthetic dataset than the sigmoid non-linearity. In fact, it remains stagnant after the first epoch.



```
train_mlp x
accuracy: 0.9972
running loss: 0.02063
epoch 016
accuracy: 0.9973
running loss: 0.02006
epoch 017
accuracy: 0.9974
running loss: 0.01954
epoch 018
accuracy: 0.9975
running loss: 0.01907
epoch 019
accuracy: 0.9976
running loss: 0.01864
Process finished with exit code 0
```

Figure 16: Training accuracy and loss with Sigmoid non-linearity.



```
train_mlp x
accuracy: 0.5409
running loss: 0.691
epoch 016
accuracy: 0.5409
running loss: 0.691
epoch 017
accuracy: 0.5409
running loss: 0.691
epoch 018
accuracy: 0.5409
running loss: 0.691
epoch 019
accuracy: 0.5409
running loss: 0.691
Process finished with exit code 0
```

Figure 17: Training accuracy and loss with Relu non-linearity.